

Information about shell project

I. Introduction

=====

For this project you will implement a simplified shell. We have provided a command parser that reads a command provided at the command prompt and creates a tree for you. Your assignment is to process that tree so the command represented by the tree is executed. The function you need to implement can be found in the `executor.c` file:

```
int execute(struct tree *t) {  
}
```

Although you need to provide a makefile for the project you can start implementing your code without it. To see what the parser does so far, compile the project code distribution as follows:

```
gcc *.c -lreadline
```

Once you compile, execute `a.out`. You will see the shell prompt (`d8sh%`). If you type any command and press enter, you will see the message "You must implement me :)". To stop the shell (which as you can see is a C program) press CTRL-C for now. Edit the `execute` function above so the function calls `print_tree(t)` to print the tree that represents the command that needs to be executed. After adding the `print_tree` function call, compile the code and execute `a.out`. Now type the `ls` command at the `d8sh%` prompt. You will see the following:

NONE: ls, IR: (null), OR: (null)

The command tree is a binary tree where a C structure represents a tree node (see the file `command.h`). Each node has a type. For this project you will implement commands associated with the NONE, AND, PIPE and SUBSHELL nodes. Each node has a left and right subtree (which are null if not subtree is present). In addition, there is an each node has an `argv` array that has the command arguments. The input field is a string that represents a file used for input redirection and the output field represents

the file used for output redirection. At this point, execute each of the commands below at the d8sh% prompt and look at the output. This will help you understand how the fields of a NONE node (also known as NONE conjunction) are populated with data. For example, t->argv[0] has the command to execute, and the rest of t->argv will have command line arguments. t->input is the file used for input redirection; t->output the file used for output redirection.

Commands to execute at the d8sh% prompt:

```
a.out
a.out < data.txt
a.out > results.txt
a.out < data.txt > results.txt
a.out arg1 arg2
cd
exit
cd /tmp
```

A NONE node is just a leaf of the tree. The NONE conjunction is the one you need to implement first. It is the one associated with cd, exit, commands with arguments (e.g., a.out arg1 arg2) and commands that use input and output redirection. This part is similar to Shell Jr. Shell commands (cd, exit) don't require fork/exec; linux commands will require fork and the child will perform an exec (execvp(t->argv[0], t->argv);) to execute the command. BEFORE the child performs the execvp call, the child needs to handle input/output redirection using dup2. If t->input is different than NULL, then open the file t->input and use dup2 to change the standard input of the child to use that file. If t->output is different than NULL, then open the file t->output and use dup2 to change the standard output of the child to that file.

Remember that in linux, a successful command execution is represented by a program returning a value of 0; any other value indicates the command failed for some reason. In tcsh you can find out the status of the last executed command by using echo \$? . The following is an example of executing a valid and invalid command in the grace cluster and the value printed after the command execution.

```
% date
Fri Jul 23 14:37:52 EDT 2021
% echo $?
0
% cat bla
cat: bla: No such file or directory
% echo $?
1
%
```

The execute function (int execute(struct tree *t)) you need to implement needs to return 0 to indicate successful execution of the command represented by tree and a non-zero value otherwise.

II. Processing commands that involve more than one node

=====

If you have a command that is just a NONE node (a root node and nothing else) you will be done after implementing the code for the NONE node, but commands can be more complex. In a linux shell we can use the following options:

1. AND && (two &)

You can connect two or more linux command with &&. Once a command fails, the rest are not executed. For example executing in the grace cluster (assuming the file lexer.h is present) will allow the following date command to execute:

```
% ls lexer.h && date
lexer.h
Fri Jul 23 14:54:11 EDT 2021
%
```

If we try the following command, where the file does not exist, the date command will not be executed.

```
% ls lexer.h && date
lexer.h
Fri Jul 23 14:54:11 EDT 2021
%|
```

Your shell will process `&&`. Execute in the `d8sh%` the following command:

```
a.out && a.out2
```

You will see the following output that is a tree with 3 nodes (an AND node (root) and two leaf nodes).

```
NONE: a.out, IR: (null), OR: (null)
```

```
&&, IR: (null), OR: (null)
```

```
NONE: a.out2, IR: (null), OR: (null)
```

2. PIPE | (single |)

A pipe allow us to connect the output of one process to the input of another. For example, the `wc` command in linux provides information about the number of lines, words in a file. Here is an example:

```
% cat info.txt
This is
a cat
and a cat
is a nice cat
% wc info.txt
4 11 38 info.txt
%
```

What if you would like to count lines and words associated with the output of `ls`? You can use a pipe where the output of `ls` will be fed to `wc` as follows:

```
% ls | wc
13 13 145
%
```

You can add more pipes (e.g., the output of `wc` can be used by another command).

Your shell will process `|`. Execute in the `d8sh%` the following command:

```
a.out | a.out2
```

You will see the following output that is a tree with 3 nodes (a PIPE node (root) and two leaf nodes).

```
NONE: a.out, IR: (null), OR: (null)
|, IR: (null), OR: (null)
NONE: a.out2, IR: (null), OR: (null)
```

To process the PIPE you will need to create a pipe (as shown in lecture) and pass the appropriate file descriptors to the children.

As you can see, the tree can be large and can have a combination of PIPE and/or AND nodes. AT THE LEAF LEVEL OF THE TREE YOU WILL FIND THE NONE nodes.

III. Processing the tree =====

At first, it looks like processing the tree is difficult due to the different types of nodes, the input/output redirection, and the different types of commands, but that is not the case. The tree has been created for you, what simplifies matters a lot! You can implement the project with simple commands that use && and | and input and output redirection. IMPORTANT: The impact of the PIPE and AND nodes in the tree is to define which file descriptors leaf nodes (NONE nodes) will use for data input and for data output.

To process the tree you need to start at the root, processing each node. Each tree node will receive from its parent node an input file descriptor (integer) that represents what the node should use as the source of input (e.g., a file, standard input, a pipe read end) and what should use for output (e.g., a file, standard output, a pipe write end). The processing that takes place at each node will depend on the kind of node, the file descriptors it receives from the parent, and whether the node has files for input/output redirection (i.e., whether t->input and t->output are different than NULL). For example, a PIPE node will create a pipe and then decide what input and output file descriptors need to send to its children. The AND node will just pass the file descriptors it receives to its children. How can the file descriptors be passed to children? There are different approaches, but

one approach is to define an auxiliary function that we can call recursively to process the tree. This function could be:

```
static int execute_aux(struct tree *t, int p_input_fd, int p_output_fd) {  
}
```

The function receives a tree, a parent input file descriptor (p_input_fd) and a parent output file descriptor (p_output_fd).

The execute function will call the above auxiliary as follows:

```
int execute(struct tree *t) {  
    return execute_aux(t, STDIN_FILENO, STDOUT_FILENO);  
}
```

As you can see the root node will receive standard input as the input file descriptor and standard output as the output file descriptor, as the root has no parent.

Every node (NONE, AND, PIPE, SUBSHELL) relies on four pieces of information to make a decision regarding that file descriptors it will pass to its children. Assuming t points to a node, the four pieces of information are:

- a. t->input (file provided for input redirection)
- b. t->output (file provided for output redirection)
- c. Input file descriptor provided by the parent node (p_input_fd)
- d. Output file descriptor provided by the parent node (p_output_fd)

The processing for each kind of node is:

1. NONE node =====

If the command is exit or cd, the shell can process it directly without doing any fork/exec/dup2. For a command other than exit or cd you need to:

- a. Fork so the child processes the command and the shell waits for the command to be executed by the child. The shell will return 0 if the child processed the command successfully and another value otherwise (any value should be fine). This value is what we refer to as the status value (the integer value returned

by the `execute` function and the `execute_aux()` function).

- b. In the child, BEFORE we execute `execvp`, we need to decide what the program we are executing will use for input/output. Does it relies on the files it may have for input/output redirection?

For a NONE node you will check whether the command relies on input/output redirection. If `t->input` is present (value other than NULL), you will ignore the input file descriptor provided by the parent (`p_input_fd`) and use the file descriptor associated with opening the file `t->input`. If `t->input` is NULL, use the file descriptor provided by the parent (`p_input_fd`). A similar processing applies to `t->output`. After deciding which file descriptors you need to use, USE `dup2` TO MAP the file descriptor 0 of the child process to the file/pipe end that will be used for input, and file descriptor 1 of the child process to the file/pipe end that will be used for output. Once this process has taken place, then `execvp` can be called and the program will execute using the correct data source/data destination. IF YOU execute `dup2` BEFORE YOU FORK, YOU WILL MESS UP THE FILE DESCRIPTORS OF THE SHELL PROCESS; THAT IS WRONG.

2. AND node

=====

- a. Process the left subtree (`t->left`) using the parent input/output file descriptors (`p_input_fd`, `p_output_fd`) that the node received.
- b. If the left subtree execution is successful, the right subtree (`t->right`) will be processed using the parent's input/output file descriptors (`p_input_fd`, `p_output_fd`). In this case the value returned by the AND node processing will be the status value returned by processing the right subtree.

If the left subtree execution failed, then the right subtree will not be processed and the AND node processing will return the status value returned by the processing of the left subtree.

3. PIPE node

=====

- a. Create a pipe
- b. One process will take care of left subtree and another of the right tree; which one is up to you.

c. For the left subtree

If `t->input` is defined, open the file and use it as input file descriptor for the left subtree; otherwise use what the parent provides (`p_input_fd`). The output file descriptor will be the write end of the pipe.

d. For the right tree

If `t->output` is defined, open the file and use it as output file descriptor for the right subtree; otherwise use what the parent provides (`p_output_fd`). The input file descriptor will be the read end of the pipe.

4. SUBSHELL node

=====

Note: A subshell is a child process launched by a shell. Parentheses start a subshell. Commands included in the parentheses will not affect the parent shell's environment (e.g., the working directory). For example, executing:

```
( cd /tmp && ls )
```

will list the files of the directory `/tmp`, but the working directory after the above command has been executed will not be `/tmp`. If you remove the parenthesis, the working directory will be `/tmp`. You can determine the working directory by executing the `pwd` command.

- a. The processing of a subshell will be done by a child process (i.e., you need to `fork()`) that processes the `t->left` subtree of the SUBSHELL node. WARNING: You will ignore the `t->right` subtree.
- b. If `t->input` is defined, open the file and use the file descriptor as the input file descriptor to process the `t->left` tree.
Notice that the original shell received standard input/output; you are doing a similar process, but this time what you are passing to the subshell is either the file descriptor received from the parent or the one associated with `t->input` (if `t->input` is different than `NULL`).
- c. A process similar to the one described in b. applies to `t->output`.

- d. The status value to return will be status of processing the left subtree.
- e. The parent will reap the child that is executing the subshell.

IV. Miscellanenous

=====

1. WARNING: Make sure your code compiles without warnings, otherwise you will lose credit. For example, make sure you comment out the `print_tree` function provided with `executor.c` once you have no longer use it. If you do not comment it out, you will lose credit.
2. WARNING: Make sure your `executor.c` file has your student id at the top of the file, otherwise you will lose credit.
3. WARNING. A `dup2` system call should only happen after a `fork()`. A common error is to perform a `dup2` before a `fork`, what changes the shell standard input/output file descriptors instead of standard input/ouput file descriptors of the child process. The command is executed, but the shell now has its file descriptors corrupted.
4. What is ambiguous input redirect and ambiguous output redirect??

Here are examples:

Ambiguous output redirect:

```
a.out > result | gone
```

Ambiguous input redirect:

```
a.out | gone < data
```

You need to recognize the above cases and stop processing and return a value other than 0. The output messages to use are:

```
printf("Ambiguous output redirect.\n");
```

```
printf("Ambiguous input redirect.\n");
```

Once you recognized the PIPE conjunction, but before you do any

processing, verify whether you have ambiguous output redirect and ambiguous input redirect. First check for ambiguous output redirect; if you identify this condition, stop processing the pipe conjunction. If there is no ambiguous output redirect, check for ambiguous input redirect and stop processing the pipe conjunction if you identify this condition.

No processing should be performed if ambiguous redirects are identified.

IMPORTANT IMPORTANT

Make sure you use the printf statements we left in these notes as students often fail tests due to typos in the expected messages.

5. WARNING: If you are adding printf statements for debugging, make sure you use fflush(stdout) so your output is guaranteed to take place.
6. In this project we don't have background processing. The loop we use to reap (in one of our lecture examples) is not needed for this project.
7. WARNING: Remember, that you create a pipe and THEN fork (otherwise the pipe will not be shared). Remember: PFChang (Pipe and Fork)

V. Videos

=====

1. The following video covers some aspects of the shell project:

<https://umd.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=10c89d23-a806-454d-955c-aa940129a742>

There is a correction that needs to be made:

In the video it is mentioned that for the && conjunction you check the t->input and t->output and open files and pass file descriptors, but for this version of the shell this is not necessary. The && conjunction will just use the file descriptors it receives from the parent and pass them to the children. The left child will be executed first and if successful the right child will be executed. There is no need to fork while processing the actual && conjunction.

2. The beginning of the following video goes into detail on how to implement subshell:

<https://umd.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=de9726ca-014d-46ea-824e-aa9601296f22#>